

Spring AI FAQ

Q: What is Spring AI?

A: Framework integrating LLMs into Spring apps. Supports ChatClient, models, embeddings, vector stores, RAG, tool calling.

Q: Core modules

A: ChatModel, EmbeddingModel, ImageModel, VectorStore, Document Readers, Advisors.

Q: ChatClient

A: Fluent API for prompts, system messages, user messages, structured output.

Q: RAG

A: Load documents, create embeddings, store vectors, retrieve relevant chunks, augment prompt.

Q: Vector stores

A: PGVector, Chroma, Milvus, Elasticsearch, Pinecone, Redis, others.

Q: Document readers

A: PDF, Text, Markdown, HTML, JSON.

Q: Embeddings

A: Convert text into vectors for semantic search.

Q: Tool calling

A: LLM invokes Java functions annotated/registered by app.

Q: Memory

A: Conversation history via chat memory implementations.

Q: Supported providers

A: OpenAI, Azure OpenAI, Ollama, Anthropic, Google Gemini, Mistral, Bedrock and more.

Q: Typical flow

A: Read docs -> Chunk -> Embed -> Store -> Retrieve -> Prompt -> LLM response.

Q: Best practices

A: Chunk 500-1000 tokens, overlap 50-100, cache embeddings, filter metadata, monitor token usage.

Q: Example use cases

A: Internal chatbot, document Q&A, code assistant, customer support, knowledge search.